

https://colab.research.google.com/drive/1ReyPTUm9_YFte7dQiaSKaxE2mKK5fRy8#scrollTo=olheeawyl5Pe

Naive Bayes

////code////////

```
# -*- coding: utf-8 -*-
"""NaiveBayes.ipynb

Automatically generated by Colab.

Original file is located at
https://colab.research.google.com/drive/1ReyPTUm9_YFte7dQiaSKaxE2mKK5fRy8
"""

import pandas as pd
import numpy as np
import re

test_data = pd.read_csv('test.csv')
df = pd.read_csv('train.csv')

print(test_data.shape)
print(df.shape)

fake_counts = test_data['fake'].value_counts()
print("Number of fake profiles:", fake_counts[1])
print("Number of real profiles:", fake_counts[0])

fake_count_training = df['fake'].value_counts()
print("Number of fake profiles:", fake_count_training[1])
print("Number of real profiles:", fake_count_training[0])

df.isnull().sum()

df.tail(5)

"""Features"""
```

```

df.columns

feature_cols = ['profile pic', 'nums/length username', 'fullname words',
                'nums/length fullname', 'name==username', 'description length',
                'external URL', 'private', '#posts', '#followers', '#follows']
X_train = df[feature_cols]

X_test = test_data[feature_cols]

y_test = test_data['fake']

y_train = df['fake']

y_train

"""#Model Training"""

# Commented out IPython magic to ensure Python compatibility.
import joblib
import matplotlib.pyplot as plt
# %matplotlib inline

from sklearn import impute
from sklearn import model_selection
from sklearn import metrics
from sklearn import preprocessing
from sklearn.metrics import roc_curve, auc, accuracy_score,
classification_report, confusion_matrix, accuracy_score, make_scorer
from sklearn.model_selection import StratifiedKFold, train_test_split,
GridSearchCV, learning_curve
from sklearn.naive_bayes import GaussianNB

def plot_learning_curve(estimator, title, X, y, ylim=None, cv=None,
                        n_jobs=1, train_sizes=np.linspace(.1, 1.0, 5)):
    """
    Generate a learning curve plot for a machine learning model.

    Parameters:
    - estimator: The machine learning model to evaluate.
    - title: The title of the plot.

```

- X: The input features.
- y: The target variable.
- ylim: Tuple, optional, the y-axis limits.
- cv: Cross-validation strategy. None by default.
- n_jobs: Number of jobs to run in parallel for cross-validation.
- train_sizes: Array of training set sizes.

Returns:

- Matplotlib plot object.
- """

```
plt.figure()
```

```
plt.title(title)
```

```
if ylim is not None:
```

```
    plt.ylim(*ylim)
```

```
plt.xlabel("Training examples")
```

```
plt.ylabel("Score")
```

```
# Generate learning curves using the learning_curve function
```

```
train_sizes, train_scores, test_scores = learning_curve(  
    estimator, X, y, cv=cv, n_jobs=n_jobs, train_sizes=train_sizes)
```

```
# Calculate mean and standard deviation of training scores
```

```
train_scores_mean = np.mean(train_scores, axis=1)
```

```
train_scores_std = np.std(train_scores, axis=1)
```

```
# Calculate mean and standard deviation of test scores
```

```
test_scores_mean = np.mean(test_scores, axis=1)
```

```
test_scores_std = np.std(test_scores, axis=1)
```

```
plt.grid()
```

```
plt.fill_between(train_sizes, train_scores_mean - train_scores_std,  
                 train_scores_mean + train_scores_std, alpha=0.1,  
                 color="r")
```

```
plt.fill_between(train_sizes, test_scores_mean - test_scores_std,  
                 test_scores_mean + test_scores_std, alpha=0.1,
```

```
color="g")
```

```

plt.plot(train_sizes, train_scores_mean, 'o-', color="r",
         label="Training score")
plt.plot(train_sizes, test_scores_mean, 'o-', color="g",
         label="Cross-validation score")

plt.legend(loc="best")
return plt

def train(X_train,y_train,X_test):
    """ Trains and predicts dataset with a Gaussian Naive Bayes """

    gnb = GaussianNB()
    gnb.fit(X_train,y_train)
    print("The best classifier is: ",gnb)
    # Estimate score
    scores = model_selection.cross_val_score(gnb, X_train,y_train, cv=5)
    print (scores)
    print('Estimated score: %0.5f (+/- %0.5f)' % (scores.mean(),
scores.std() / 2))
    title = 'Learning Curves (Naive Bayes)'
    plot_learning_curve(gnb, title, X_train, y_train, cv=5)
    plt.show()
    # Predict
    y_pred = gnb.predict(X_test)
    return y_test,y_pred

# Start training models
y_test, y_pred = train(X_train, y_train, X_test)

"""Plot Confusion Matrix"""

def plot_confusion_matrix(cm, title='Confusion matrix Naive Bayes',
cmap=plt.cm.Blues):
    target_names=['Fake','Genuine(Real)']
    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(target_names))
    plt.xticks(tick_marks, target_names)
    plt.yticks(tick_marks, target_names)

```

```

plt.tight_layout()
plt.ylabel('True label')
plt.xlabel('Predicted label')

cm=confusion_matrix(y_test, y_pred)

print('Confusion Matrix[without Normalization]')
print(cm)

plot_confusion_matrix(cm)

print(classification_report(y_test, y_pred))

```

<https://colab.research.google.com/drive/1ljaM969ZpETrPZw4B8K5sjYnWU85M3c8>

Multi Layer Perceptron

```

# -*- coding: utf-8 -*-
"""Fake Instagram Profile Detection Model ....

Automatically generated by Colab.

Original file is located at
https://colab.research.google.com/drive/1ljaM969ZpETrPZw4B8K5sjYnWU85M3c8
"""

import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns

import tensorflow as tf

```

```
from tensorflow import keras
from tensorflow.keras.layers import Dense, Activation, Dropout
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.metrics import Accuracy

from sklearn import metrics
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import
classification_report, accuracy_score, roc_curve, confusion_matrix

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

# Load the training dataset
instagram_df_train=pd.read_csv('train.csv')
instagram_df_train

# Load the testing data
instagram_df_test=pd.read_csv('test.csv')
instagram_df_test

"""# Statistical Analysis"""

instagram_df_train.head()

instagram_df_train.tail()

# Getting dataframe info
instagram_df_train.info()

# Get the statistical summary of the dataframe
instagram_df_train.describe()

# Checking if null values exist
instagram_df_train.isnull().sum()

# Get the number of unique values in the "profile pic" feature
instagram_df_train['profile pic'].value_counts()
```

```

# Get the number of unique values in "fake" (Target column)
instagram_df_train['fake'].value_counts()

"""# Data Visualization"""

# Visualize the data
sns.countplot(instagram_df_train['fake'])
plt.show()

# Visualize the private column data
sns.countplot(instagram_df_train['private'])
plt.show()

# Visualize the "profile pic" column data
sns.countplot(instagram_df_train['profile pic'])
plt.show()

# Visualize the data
plt.figure(figsize = (20, 10))
sns.distplot(instagram_df_train['nums/length username'])
plt.show()

# Correlation plot
plt.figure(figsize=(20, 20))
cm = instagram_df_train.corr()
ax = plt.subplot()
sns.heatmap(cm, annot = True, ax = ax)
plt.show()

"""# Data Modelling"""

# Training and testing dataset (inputs)
X_train = instagram_df_train.drop(columns = ['fake'])
X_test = instagram_df_test.drop(columns = ['fake'])
X_train

# Training and testing dataset (Outputs)
y_train = instagram_df_train['fake']
y_test = instagram_df_test['fake']

```

```

y_train

# Scale the data before training the model
from sklearn.preprocessing import StandardScaler, MinMaxScaler

scaler_x = StandardScaler()
X_train = scaler_x.fit_transform(X_train)
X_test = scaler_x.transform(X_test)

y_train = tf.keras.utils.to_categorical(y_train, num_classes = 2)
y_test = tf.keras.utils.to_categorical(y_test, num_classes = 2)

y_train

import tensorflow.keras
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout

model = Sequential()
model.add(Dense(50, input_dim=11, activation='relu'))
model.add(Dense(150, activation='relu'))
model.add(Dropout(0.3))
model.add(Dense(150, activation='relu'))
model.add(Dropout(0.3))
model.add(Dense(25, activation='relu'))
model.add(Dropout(0.3))
model.add(Dense(2, activation='softmax'))

model.summary()

model.compile(optimizer = 'adam', loss = 'categorical_crossentropy',
metrics = ['accuracy'])

epochs_hist = model.fit(X_train, y_train, epochs = 50, verbose = 1,
validation_split = 0.1)

"""# Model Validation and Results"""

print(epochs_hist.history.keys())

```



```

plt.plot(epochs_hist.history['loss'])
plt.plot(epochs_hist.history['val_loss'])

plt.title('Model Loss Progression During Training/Validation')
plt.ylabel('Training and Validation Losses')
plt.xlabel('Epoch Number')
plt.legend(['Training Loss', 'Validation Loss'])
plt.show()

predicted = model.predict(X_test)

predicted_value = []
test = []
for i in predicted:
    predicted_value.append(np.argmax(i))

for i in y_test:
    test.append(np.argmax(i))

print(classification_report(test, predicted_value))

plt.figure(figsize=(10, 10))
cm=confusion_matrix(test, predicted_value)
sns.heatmap(cm, annot=True)
plt.show()

```

<https://colab.research.google.com/drive/16883jLs7yFhpHoWTgYYyLsOXfWfUTs0g?usp=sharing>

Random Forest

```

# -*- coding: utf-8 -*-
"""RandomForest.ipynb

Automatically generated by Colab.

```

Original file is located at

<https://colab.research.google.com/drive/16883jLs7yFhpHoWTgYYyLsOXfWfUTs0g>
"""

```
import pandas as pd
import numpy as np
import re
```

```
test_data = pd.read_csv('test.csv')
df = pd.read_csv('train.csv')
```

```
print(test_data.shape)
print(df.shape)
```

```
df.isnull().sum()
```

```
df.tail(5)
```

```
"""Features"""
```

```
df.columns
```

```
feature_cols = ['profile pic', 'nums/length username', 'fullname words',
                'nums/length fullname', 'name==username', 'description length',
                'external URL', 'private', '#posts', '#followers', '#follows']
```

```
X_train = df[feature_cols]
```

```
X_test = test_data[feature_cols]
```

```
y_test = test_data['fake']
```

```
y_train = df['fake']
```

```
y_train
```

```
"""#Model Training"""
```

```

# Commented out IPython magic to ensure Python compatibility.
import joblib
import matplotlib.pyplot as plt
# %matplotlib inline

from sklearn import impute
from sklearn import model_selection
from sklearn import metrics
from sklearn import preprocessing
from sklearn.metrics import roc_curve, auc, accuracy_score,
classification_report, confusion_matrix, accuracy_score, make_scorer
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import StratifiedKFold, train_test_split,
GridSearchCV, learning_curve
from xgboost import XGBClassifier
from sklearn.svm import SVC

def plot_learning_curve(estimator, title, X, y, ylim=None, cv=None,
                        n_jobs=1, train_sizes=np.linspace(.1, 1.0, 5)):
    """
    Generate a learning curve plot for a machine learning model.

    Parameters:
    - estimator: The machine learning model to evaluate.
    - title: The title of the plot.
    - X: The input features.
    - y: The target variable.
    - ylim: Tuple, optional, the y-axis limits.
    - cv: Cross-validation strategy. None by default.
    - n_jobs: Number of jobs to run in parallel for cross-validation.
    - train_sizes: Array of training set sizes.

    Returns:
    - Matplotlib plot object.
    """

    plt.figure()
    plt.title(title)

    if ylim is not None:

```

```

plt.ylim(*ylim)

plt.xlabel("Training examples")
plt.ylabel("Score")

# Generate learning curves using the learning_curve function
train_sizes, train_scores, test_scores = learning_curve(
    estimator, X, y, cv=cv, n_jobs=n_jobs, train_sizes=train_sizes)

# Calculate mean and standard deviation of training scores
train_scores_mean = np.mean(train_scores, axis=1)
train_scores_std = np.std(train_scores, axis=1)

# Calculate mean and standard deviation of test scores
test_scores_mean = np.mean(test_scores, axis=1)
test_scores_std = np.std(test_scores, axis=1)

plt.grid()
plt.fill_between(train_sizes, train_scores_mean - train_scores_std,
                 train_scores_mean + train_scores_std, alpha=0.1,
                 color="r")
plt.fill_between(train_sizes, test_scores_mean - test_scores_std,
                 test_scores_mean + test_scores_std, alpha=0.1,
color="g")

plt.plot(train_sizes, train_scores_mean, 'o-', color="r",
         label="Training score")
plt.plot(train_sizes, test_scores_mean, 'o-', color="g",
         label="Cross-validation score")

plt.legend(loc="best")
return plt

def train(X_train,y_train,X_test):
    """ Trains and predicts dataset with a Random Forest classifier """

    clf=RandomForestClassifier(n_estimators=40)
    clf.fit(X_train,y_train)
    print("The best classifier is: ",clf)
    # Estimate score

```

```

    scores = model_selection.cross_val_score(clf, X_train,y_train, cv=5)
    print (scores)
    print('Estimated score: %0.5f (+/- %0.5f)' % (scores.mean(),
scores.std() / 2))
    title = 'Learning Curves (Random Forest)'
    plot_learning_curve(clf, title, X_train, y_train, cv=5)
    plt.show()
    # Predict
    y_pred = clf.predict(X_test)
    # print("output")
    # print(clf.predict([[100,872,924,20,15,2,3]]))
    # print(clf.predict([[1234,15,104,1150,0,2,1]]))
    # print(clf.predict([[60,14,592,0,0,-2,1]]))
    # print(clf.predict([[16935,2242,918,44,124,-2,1]]))
    # print(clf.predict([[1294,254,276,89,8,2,5]]))
    # print(clf.predict([[44,21,614,0,0,-2,1]]))#fake
    # print(clf.predict([[1386,88,267,106,0,2,1]]))#g
    # print(clf.predict([[4149,279,64,20,22,-1,1]]))#g
    # print(clf.predict([[732,191,656,1,1,1,3]]))#g
    return y_test,y_pred

# Start training models
y_test, y_pred = train(X_train, y_train, X_test)

"""Plot Confusion Matrix"""

def plot_confusion_matrix(cm, title='Confusion matrix',
cmap=plt.cm.Blues):
    target_names=['Fake','Genuine(Real)']
    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(target_names))
    plt.xticks(tick_marks, target_names)
    plt.yticks(tick_marks, target_names)
    plt.tight_layout()
    plt.ylabel('True label')
    plt.xlabel('Predicted label')

cm=confusion_matrix(y_test, y_pred)

```

```
print('Confusion Matrix[without Normalization]')
print(cm)

plot_confusion_matrix(cm)

print(classification_report(y_test,y_pred))
```

<https://colab.research.google.com/drive/1Yv9oblxdbrsah0nmZriM4YCW7LRbjsSP?usp=sharing>
SVM

```
# -*- coding: utf-8 -*-
"""SVM.ipynb

Automatically generated by Colab.

Original file is located at
https://colab.research.google.com/drive/1Yv9oblxdbrsah0nmZriM4YCW7LRbjsSP
"""

import pandas as pd
import numpy as np
import re

test_data = pd.read_csv('test.csv')
df = pd.read_csv('train.csv')

print(test_data.shape)
print(df.shape)

df.isnull().sum()
```

```

df.tail(5)

"""Features"""

df.columns

feature_cols = ['profile pic', 'nums/length username', 'fullname words',
                'nums/length fullname', 'name==username', 'description length',
                'external URL', 'private', '#posts', '#followers', '#follows']
X_train = df[feature_cols]

X_test = test_data[feature_cols]

y_test = test_data['fake']


y_train = df['fake']

y_train

"""#Model Training"""

# Commented out IPython magic to ensure Python compatibility.
import joblib
import matplotlib.pyplot as plt
# %matplotlib inline

from sklearn import impute
from sklearn import model_selection
from sklearn import metrics
from sklearn import preprocessing
from sklearn.metrics import roc_curve, auc, accuracy_score,
classification_report, confusion_matrix, accuracy_score, make_scorer
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import StratifiedKFold, train_test_split,
GridSearchCV, learning_curve
from xgboost import XGBClassifier
from sklearn.svm import SVC

```

```

def plot_learning_curve(estimator, title, X, y, ylim=None, cv=None,
                        n_jobs=1, train_sizes=np.linspace(.1, 1.0, 5)):
    """
    Generate a learning curve plot for a machine learning model.

    Parameters:
    - estimator: The machine learning model to evaluate.
    - title: The title of the plot.
    - X: The input features.
    - y: The target variable.
    - ylim: Tuple, optional, the y-axis limits.
    - cv: Cross-validation strategy. None by default.
    - n_jobs: Number of jobs to run in parallel for cross-validation.
    - train_sizes: Array of training set sizes.

    Returns:
    - Matplotlib plot object.
    """

    plt.figure()
    plt.title(title)

    if ylim is not None:
        plt.ylim(*ylim)

    plt.xlabel("Training examples")
    plt.ylabel("Score")

    # Generate learning curves using the learning_curve function
    train_sizes, train_scores, test_scores = learning_curve(
        estimator, X, y, cv=cv, n_jobs=n_jobs, train_sizes=train_sizes)

    # Calculate mean and standard deviation of training scores
    train_scores_mean = np.mean(train_scores, axis=1)
    train_scores_std = np.std(train_scores, axis=1)

    # Calculate mean and standard deviation of test scores
    test_scores_mean = np.mean(test_scores, axis=1)
    test_scores_std = np.std(test_scores, axis=1)

```



```

plt.grid()
plt.fill_between(train_sizes, train_scores_mean - train_scores_std,
                 train_scores_mean + train_scores_std, alpha=0.1,
                 color="r")
plt.fill_between(train_sizes, test_scores_mean - test_scores_std,
                 test_scores_mean + test_scores_std, alpha=0.1,
color="g")

plt.plot(train_sizes, train_scores_mean, 'o-', color="r",
         label="Training score")
plt.plot(train_sizes, test_scores_mean, 'o-', color="g",
         label="Cross-validation score")

plt.legend(loc="best")
return plt

def train(X_train, y_train, X_test, y_test):
    filename = 'model.ckpt'
    param_grid = {
        'C': [0.1, 1, 10],
        'gamma': [0.01, 0.1, 1],
        'kernel': ['rbf']
    }

    svm_clf = SVC()
    scorer = make_scorer(accuracy_score)
    grid_search = GridSearchCV(svm_clf, param_grid, cv=5, scoring=scorer,
n_jobs=-1)
    grid_search.fit(X_train, y_train)

    best_svm_model = grid_search.best_estimator_
    joblib.dump(best_svm_model, filename)

    print("Training Datasets using SVC")
    print("Best Hyperparameters:", grid_search.best_params_)
    y_pred = best_svm_model.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)
    print("Accuracy of the Best Model:", accuracy)

    title = 'Learning Curves (SVM)'

```

9

```
plot_learning_curve(best_svm_model, title, X_train, y_train, cv=5)
plt.show()
# print("output svm ")
# print(best_svm_model.predict([[-2,16935,2242,918,44,124,1]]))#real
# print(best_svm_model.predict([[2,1234,15,104,1150,0,1]]))#fake
# print(best_svm_model.predict([[-2,60,14,592,0,0,1]]))#fake
# print(best_svm_model.predict([[2,1294,254,276,89,8,5]]))#real
# print(best_svm_model.predict([[2,20370,5470,2385,145,52,5]]))#real
# print(best_svm_model.predict([[-2,3131,506,381,9,40,1]]))#real
# print(best_svm_model.predict([[0,4024,264,87,323,16,1]]))#real
# print(best_svm_model.predict([[0,40586,640,622,1118,32,1]]))#real
# print(best_svm_model.predict([[2,2016,62,64,13,0,5]]))#real
# print(best_svm_model.predict([[2,2016,62,64,13,0,5]]))#real

return y_test, y_pred

# Start training models
y_test, y_pred = train(X_train, y_train, X_test, y_test)

"""Plot Confusion Matrix"""

def plot_confusion_matrix(cm, title='Confusion matrix',
cmap=plt.cm.Blues):
    target_names=['Fake','Genuine(Real)']
    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(target_names))
    plt.xticks(tick_marks, target_names)
    plt.yticks(tick_marks, target_names)
    plt.tight_layout()
    plt.ylabel('True label')
    plt.xlabel('Predicted label')

cm=confusion_matrix(y_test, y_pred)

print('Confusion Matrix[without Normalization]')
print(cm)
```

```
plot_confusion_matrix(cm)

print(classification_report(y_test,y_pred))
```

<https://colab.research.google.com/drive/1jC1Q7aSg5ctlvhFZivAkwgwhACuzgDGE?usp=sharing>
KNN

```
# -*- coding: utf-8 -*-
"""KNN.ipynb

Automatically generated by Colab.

Original file is located at
https://colab.research.google.com/drive/1jC1Q7aSg5ctlvhFZivAkwgwhACuzgDGE
"""

import pandas as pd
import numpy as np
import re

test_data = pd.read_csv('test.csv')
df = pd.read_csv('train.csv')

print(test_data.shape)
print(df.shape)

df.isnull().sum()

df.tail(5)

"""Features"""
```

```

df.columns

feature_cols = ['profile pic', 'nums/length username', 'fullname words',
                'nums/length fullname', 'name==username', 'description length',
                'external URL', 'private', '#posts', '#followers', '#follows']
X_train = df[feature_cols]

X_test = test_data[feature_cols]

y_test = test_data['fake']

y_train = df['fake']

y_train

"""#Model Training"""

# Commented out IPython magic to ensure Python compatibility.
import joblib
import matplotlib.pyplot as plt
# %matplotlib inline

from sklearn import impute
from sklearn import model_selection
from sklearn import metrics
from sklearn import preprocessing
from sklearn.metrics import roc_curve, auc, accuracy_score,
classification_report, confusion_matrix, accuracy_score, make_scorer
from sklearn.model_selection import StratifiedKFold, train_test_split,
GridSearchCV, learning_curve
from sklearn.neighbors import KNeighborsClassifier

def plot_learning_curve(estimator, title, X, y, ylim=None, cv=None,
                        n_jobs=1, train_sizes=np.linspace(.1, 1.0, 5)):
    """
    Generate a learning curve plot for a machine learning model.

    Parameters:
    - estimator: The machine learning model to evaluate.

```

- title: The title of the plot.
- X: The input features.
- y: The target variable.
- ylim: Tuple, optional, the y-axis limits.
- cv: Cross-validation strategy. None by default.
- n_jobs: Number of jobs to run in parallel for cross-validation.
- train_sizes: Array of training set sizes.

Returns:

- Matplotlib plot object.

"""

```
plt.figure()
```

```
plt.title(title)
```

```
if ylim is not None:
```

```
    plt.ylim(*ylim)
```

```
plt.xlabel("Training examples")
```

```
plt.ylabel("Score")
```

```
# Generate learning curves using the learning_curve function
```

```
train_sizes, train_scores, test_scores = learning_curve(  
    estimator, X, y, cv=cv, n_jobs=n_jobs, train_sizes=train_sizes)
```

```
# Calculate mean and standard deviation of training scores
```

```
train_scores_mean = np.mean(train_scores, axis=1)
```

```
train_scores_std = np.std(train_scores, axis=1)
```

```
# Calculate mean and standard deviation of test scores
```

```
test_scores_mean = np.mean(test_scores, axis=1)
```

```
test_scores_std = np.std(test_scores, axis=1)
```

```
plt.grid()
```

```
plt.fill_between(train_sizes, train_scores_mean - train_scores_std,  
                 train_scores_mean + train_scores_std, alpha=0.1,  
                 color="r")
```

```
plt.fill_between(train_sizes, test_scores_mean - test_scores_std,  
                 test_scores_mean + test_scores_std, alpha=0.1,
```

```
color="g")
```

```

plt.plot(train_sizes, train_scores_mean, 'o-', color="r",
         label="Training score")
plt.plot(train_sizes, test_scores_mean, 'o-', color="g",
         label="Cross-validation score")

plt.legend(loc="best")
return plt

def train(X_train,y_train,X_test):
    """ Trains and predicts dataset with a K Neighbours """

    knn = KNeighborsClassifier(n_neighbors=3)
    knn.fit(X_train,y_train)
    print("The best classifier is: ",knn)
    # Estimate score
    scores = model_selection.cross_val_score(knn, X_train,y_train, cv=5)
    print (scores)
    print('Estimated score: %0.5f (+/- %0.5f)' % (scores.mean(),
scores.std() / 2))
    title = 'Learning Curves (K Neighbours)'
    plot_learning_curve(knn, title, X_train, y_train, cv=5)
    plt.show()
    # Predict
    y_pred = knn.predict(X_test)
    # print("output")
    # print(clf.predict([[100,872,924,20,15,2,3]]))
    # print(clf.predict([[1234,15,104,1150,0,2,1]]))
    # print(clf.predict([[60,14,592,0,0,-2,1]]))
    # print(clf.predict([[16935,2242,918,44,124,-2,1]]))
    # print(clf.predict([[1294,254,276,89,8,2,5]]))
    # print(clf.predict([[44,21,614,0,0,-2,1]]))#fake
    # print(clf.predict([[1386,88,267,106,0,2,1]]))#g
    # print(clf.predict([[4149,279,64,20,22,-1,1]]))#g
    # print(clf.predict([[732,191,656,1,1,1,3]]))#g
    return y_test,y_pred

# Start training models
y_test, y_pred = train(X_train, y_train, X_test)

```

```

"""Plot Confusion Matrix"""

def plot_confusion_matrix(cm, title='Confusion matrix',
cmap=plt.cm.Blues):
    target_names=['Fake', 'Genuine(Real)']
    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(target_names))
    plt.xticks(tick_marks, target_names)
    plt.yticks(tick_marks, target_names)
    plt.tight_layout()
    plt.ylabel('True label')
    plt.xlabel('Predicted label')

cm=confusion_matrix(y_test, y_pred)

print('Confusion Matrix[without Normalization]')
print(cm)

plot_confusion_matrix(cm)

print(classification_report(y_test,y_pred))

```

<https://colab.research.google.com/drive/16zIT7ljM-SquoQFEVgRabOP7G4LInRD7?usp=sharing>

Logistic Regression

```

# -*- coding: utf-8 -*-

"""LogisticRegression.ipynb

```

Automatically generated by Colab.

Original file is located at

<https://colab.research.google.com/drive/16zIT7IjM-SguoQFEVgRabOP7G4LlnRD7>

"""

```
import pandas as pd
```

```
import numpy as np
```

```
import re
```

```
test_data = pd.read_csv('test.csv')
```

```
df = pd.read_csv('train.csv')
```

```
print(test_data.shape)
```

```
print(df.shape)
```

```
df.isnull().sum()
```

```
df.tail(5)
```

"""Features"""

```
df.columns
```

```
feature_cols = ['profile pic', 'nums/length username', 'fullname words',  
                'nums/length fullname', 'name==username', 'description length',  
                'external URL', 'private', '#posts', '#followers', '#follows']
```

```
X_train = df[feature_cols]
```

```
X_test = test_data[feature_cols]
```

```
y_test = test_data['fake']
```

```
y_train = df['fake']
```

```
y_train
```

"""#Model Training"""


```

# Commented out IPython magic to ensure Python compatibility.
import joblib
import matplotlib.pyplot as plt
# %matplotlib inline

from sklearn import impute
from sklearn import model_selection
from sklearn import metrics
from sklearn import preprocessing
from sklearn.metrics import roc_curve, auc, accuracy_score,
classification_report, confusion_matrix, accuracy_score, make_scorer
from sklearn.model_selection import StratifiedKFold, train_test_split,
GridSearchCV, learning_curve
from sklearn.linear_model import LogisticRegression

def plot_learning_curve(estimator, title, X, y, ylim=None, cv=None,
                        n_jobs=1, train_sizes=np.linspace(.1, 1.0, 5)):
    """
    Generate a learning curve plot for a machine learning model.

    Parameters:
    - estimator: The machine learning model to evaluate.
    - title: The title of the plot.
    - X: The input features.
    - y: The target variable.
    - ylim: Tuple, optional, the y-axis limits.
    - cv: Cross-validation strategy. None by default.
    - n_jobs: Number of jobs to run in parallel for cross-validation.
    - train_sizes: Array of training set sizes.

    Returns:
    - Matplotlib plot object.
    """

    plt.figure()
    plt.title(title)

    if ylim is not None:
        plt.ylim(*ylim)

```

```

plt.xlabel("Training examples")
plt.ylabel("Score")

# Generate learning curves using the learning_curve function
train_sizes, train_scores, test_scores = learning_curve(
    estimator, X, y, cv=cv, n_jobs=n_jobs, train_sizes=train_sizes)

# Calculate mean and standard deviation of training scores
train_scores_mean = np.mean(train_scores, axis=1)
train_scores_std = np.std(train_scores, axis=1)

# Calculate mean and standard deviation of test scores
test_scores_mean = np.mean(test_scores, axis=1)
test_scores_std = np.std(test_scores, axis=1)

plt.grid()
plt.fill_between(train_sizes, train_scores_mean - train_scores_std,
                 train_scores_mean + train_scores_std, alpha=0.1,
                 color="r")
plt.fill_between(train_sizes, test_scores_mean - test_scores_std,
                 test_scores_mean + test_scores_std, alpha=0.1,
color="g")

plt.plot(train_sizes, train_scores_mean, 'o-', color="r",
         label="Training score")
plt.plot(train_sizes, test_scores_mean, 'o-', color="g",
         label="Cross-validation score")

plt.legend(loc="best")
return plt

def train(X_train,y_train,X_test):
    """ Trains and predicts dataset with a Logistic Regression """

    clf = LogisticRegression(random_state=0,max_iter=500)
    clf.fit(X_train,y_train)
    print("The best classifier is: ",clf)
    # Estimate score
    scores = model_selection.cross_val_score(clf, X_train,y_train, cv=5)
    print (scores)

```

```

    print('Estimated score: %0.5f (+/- %0.5f)' % (scores.mean(),
scores.std() / 2))
    title = 'Learning Curves (Logistic Regression)'
    plot_learning_curve(clf, title, X_train, y_train, cv=5)
    plt.show()
    # Predict
    y_pred = clf.predict(X_test)
    # print("output")
    # print(clf.predict([[100,872,924,20,15,2,3]]))
    # print(clf.predict([[1234,15,104,1150,0,2,1]]))
    # print(clf.predict([[60,14,592,0,0,-2,1]]))
    # print(clf.predict([[16935,2242,918,44,124,-2,1]]))
    # print(clf.predict([[1294,254,276,89,8,2,5]]))
    # print(clf.predict([[44,21,614,0,0,-2,1]])) #fake
    # print(clf.predict([[1386,88,267,106,0,2,1]])) #g
    # print(clf.predict([[4149,279,64,20,22,-1,1]])) #g
    # print(clf.predict([[732,191,656,1,1,1,3]])) #g
    return y_test, y_pred

# Start training models
y_test, y_pred = train(X_train, y_train, X_test)

"""Plot Confusion Matrix"""

def plot_confusion_matrix(cm, title='Confusion matrix',
cmap=plt.cm.Blues):
    target_names=['Fake', 'Genuine (Real)']
    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(target_names))
    plt.xticks(tick_marks, target_names)
    plt.yticks(tick_marks, target_names)
    plt.tight_layout()
    plt.ylabel('True label')
    plt.xlabel('Predicted label')

cm=confusion_matrix(y_test, y_pred)

print('Confusion Matrix[without Normalization]')

```

```
print(cm)

plot_confusion_matrix(cm)

print(classification_report(y_test,y_pred))
```

<https://colab.research.google.com/drive/1AE1qmi0ERF5DUnJfkr-POHi8T5svRBbF?usp=sharing>
Gradient Boost Classifier

```
# -*- coding: utf-8 -*-
"""GradientBoost.ipynb

Automatically generated by Colab.

Original file is located at
https://colab.research.google.com/drive/1AE1qmi0ERF5DUnJfkr-POHi8T5svRBbF
"""

import pandas as pd
import numpy as np
import re

test_data = pd.read_csv('test.csv')
df = pd.read_csv('train.csv')

print(test_data.shape)
print(df.shape)

df.isnull().sum()

df.tail(5)

"""Features"""
```

```

df.columns

feature_cols = ['profile pic', 'nums/length username', 'fullname words',
                'nums/length fullname', 'name==username', 'description length',
                'external URL', 'private', '#posts', '#followers', '#follows']
X_train = df[feature_cols]

X_test = test_data[feature_cols]

y_test = test_data['fake']

y_train = df['fake']

y_train

"""#Model Training"""

# Commented out IPython magic to ensure Python compatibility.
import joblib
import matplotlib.pyplot as plt
# %matplotlib inline

from sklearn import impute
from sklearn import model_selection
from sklearn import metrics
from sklearn import preprocessing
from sklearn.metrics import roc_curve, auc, accuracy_score,
classification_report, confusion_matrix, accuracy_score, make_scorer
from sklearn.model_selection import StratifiedKFold, train_test_split,
GridSearchCV, learning_curve
from sklearn.ensemble import GradientBoostingClassifier

def plot_learning_curve(estimator, title, X, y, ylim=None, cv=None,
                        n_jobs=1, train_sizes=np.linspace(.1, 1.0, 5)):
    """
    Generate a learning curve plot for a machine learning model.

    Parameters:
    - estimator: The machine learning model to evaluate.

```

- title: The title of the plot.
- X: The input features.
- y: The target variable.
- ylim: Tuple, optional, the y-axis limits.
- cv: Cross-validation strategy. None by default.
- n_jobs: Number of jobs to run in parallel for cross-validation.
- train_sizes: Array of training set sizes.

Returns:

- Matplotlib plot object.

"""

```
plt.figure()
```

```
plt.title(title)
```

```
if ylim is not None:
```

```
    plt.ylim(*ylim)
```

```
plt.xlabel("Training examples")
```

```
plt.ylabel("Score")
```

```
# Generate learning curves using the learning_curve function
```

```
train_sizes, train_scores, test_scores = learning_curve(  
    estimator, X, y, cv=cv, n_jobs=n_jobs, train_sizes=train_sizes)
```

```
# Calculate mean and standard deviation of training scores
```

```
train_scores_mean = np.mean(train_scores, axis=1)
```

```
train_scores_std = np.std(train_scores, axis=1)
```

```
# Calculate mean and standard deviation of test scores
```

```
test_scores_mean = np.mean(test_scores, axis=1)
```

```
test_scores_std = np.std(test_scores, axis=1)
```

```
plt.grid()
```

```
plt.fill_between(train_sizes, train_scores_mean - train_scores_std,  
                 train_scores_mean + train_scores_std, alpha=0.1,  
                 color="r")
```

```
plt.fill_between(train_sizes, test_scores_mean - test_scores_std,  
                 test_scores_mean + test_scores_std, alpha=0.1,
```

```
color="g")
```

```

plt.plot(train_sizes, train_scores_mean, 'o-', color="r",
         label="Training score")
plt.plot(train_sizes, test_scores_mean, 'o-', color="g",
         label="Cross-validation score")

plt.legend(loc="best")
return plt

def train(X_train,y_train,X_test):
    """ Trains and predicts dataset with a Gradient Boost """

    clf = GradientBoostingClassifier(n_estimators=100, learning_rate=1.0,
max_depth=1, random_state=0)
    clf.fit(X_train,y_train)
    print("The best classifier is: ",clf)
    # Estimate score
    scores = model_selection.cross_val_score(clf, X_train,y_train, cv=5)
    print (scores)
    print('Estimated score: %0.5f (+/- %0.5f)' % (scores.mean(),
scores.std() / 2))
    title = 'Learning Curves (Gradient Boost)'
    plot_learning_curve(clf, title, X_train, y_train, cv=5)
    plt.show()
    # Predict
    y_pred = clf.predict(X_test)
    return y_test,y_pred

# Start training models
y_test, y_pred = train(X_train, y_train, X_test)

"""Plot Confusion Matrix"""

def plot_confusion_matrix(cm, title='Confusion matrix',
cmap=plt.cm.Blues):
    target_names=['Fake','Genuine (Real)']
    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(target_names))

```

```
plt.xticks(tick_marks, target_names)
plt.yticks(tick_marks, target_names)
plt.tight_layout()
plt.ylabel('True label')
plt.xlabel('Predicted label')

cm=confusion_matrix(y_test, y_pred)

print('Confusion Matrix[without Normalization]')
print(cm)

plot_confusion_matrix(cm)

print(classification_report(y_test,y_pred))
```